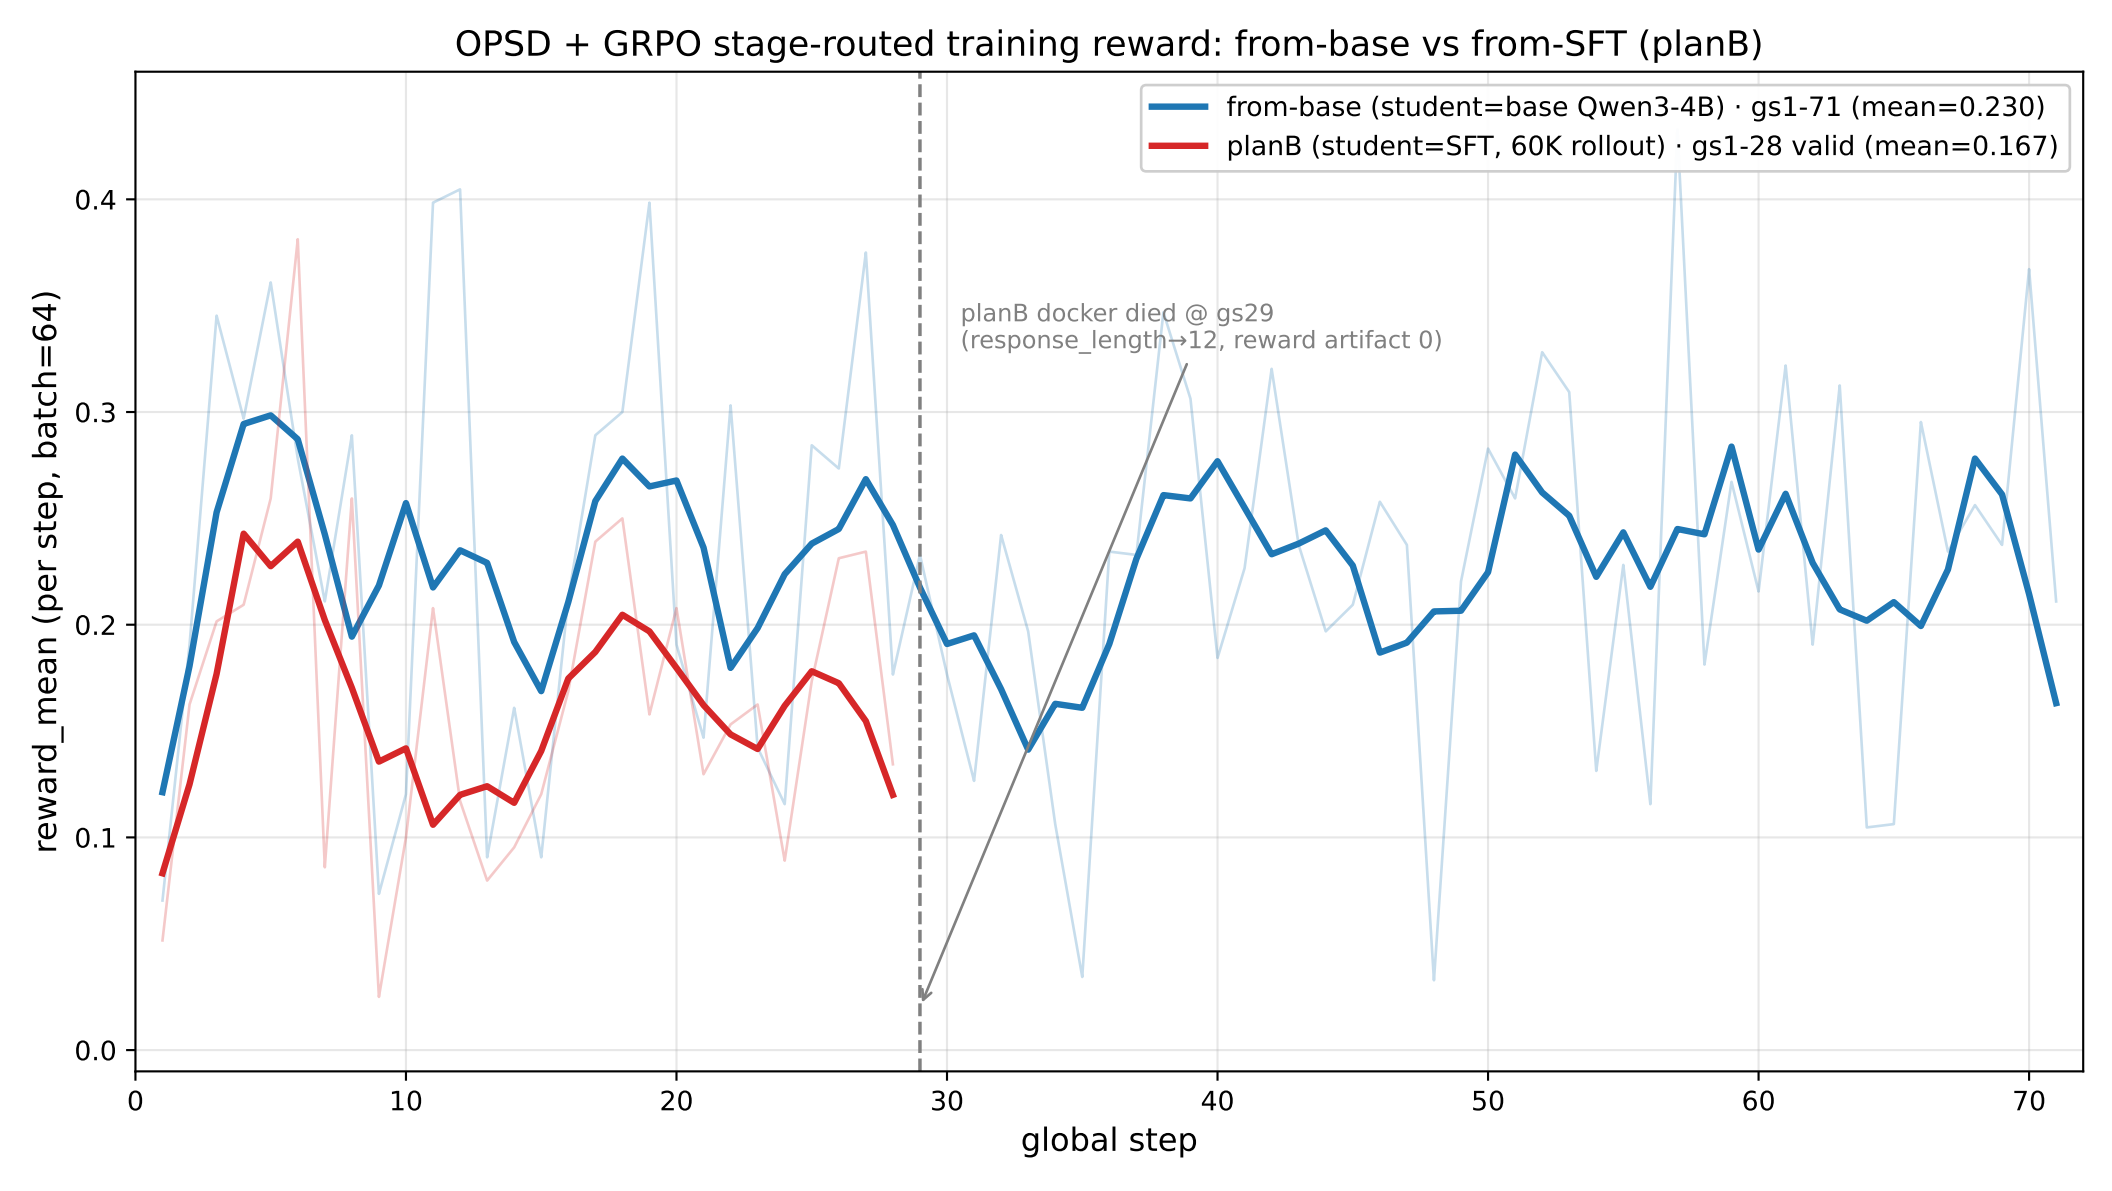
OPSD + GRPO stage-routed training reward: from-base vs from-SFT (planB)
from-base (student=base Qwen3-4B) · gs1-71 (mean=0.230)
planB (student=SFT, 60K rollout) · gs1-28 valid (mean=0.167)
planB docker died @ gs29
(response_length→12, reward artifact 0)
reward_mean (per step, batch=64)
global step
0.0
0.1
0.2
0.3
0.4
0
10
20
30
40
50
60
70